

Proposal: Replace the sourcing reverse block scan with an index lookup

Status	Proposal
Area	JetStream stream sourcing (<code>server/stream.go</code>) + store index (<code>server/filestore.go</code> , <code>server/memstore.go</code>)
Date	2026-06-06
Companion	<code>sourcing-durable-resume.md</code> (lifecycle, durable consumer, triggers)
Target	<code>startingSequenceForSources</code> (<code>stream.go:4694</code>) and <code>setStartingSequenceForSources</code> (<code>stream.go:4566</code>)


1. The problem in one paragraph

On hub-stream **leader election / restart**, `setupSourceConsumers` calls `startingSequenceForSources` **unconditionally** (`stream.go:4821`). That function reconstructs each source's last sourced sequence by walking the stream's **own store backwards from `LastSeq`**, one block at a time (`LoadPrevMsgMulti` , `filestore.go:9425`), reading the `Nats-Stream-Source` header off matching messages, until **every** source has been located.

Note: the walk is not as naive as message-by-message — within each block the per-block index (`fss`) skips non-matching messages, and the per-source sublist narrows as sources are found. But `LoadPrevMsgMulti` still **visits every block** from `LastSeq` back to the match (loading/decompressing cold ones), and the outer loop **re-walks** as the sublist narrows. So the cost grows with both the number of blocks back to the least-recently-active source (worst case the **entire store**) and the number of sources — all under `mset.mu` .

The reverse block scan: rebuild each source's last sequence from stored headers

startingSequenceForSources walks back from LastSeq, LoadPrevMsgMulti loads (+decompresses) a block at a time until every source is found.



Why it can be slow on a limits-based stream

A single quiet/sparse source (S3) forces the scan past every newer block — worst case the entire stream. Each uncached block is read from disk, decrypted and decompressed. With many sources and a large store, leader election stalls under mset.mu while this runs.

2. The key insight

The store **already maintains a per-subject index** that can return the last sequence for a subject *without* walking messages:

- `fs.psim` — per-subject info, including `lblk` = the **last block index** that contains each subject (`filestore.go:3843-3852` , `8980-8984`).
- `mb.fss` — per-block subject state with `Last` per subject (`filestore.go:9003-9015`).


This is what powers `LoadLastMsg(subject) → loadLast ( filestore.go:8939 )`: for a literal subject it jumps straight to `info.lblk` and reads that subject's last message — **one targeted block load**, regardless of how far back it is. `MultiLastSeqs ( filestore.go:3806 )` does the same for a batch of filters in a single index pass.

The current scan ignores this index and instead does a generic backward message walk. The header it needs (the **origin** sequence) is right there in the message that `LoadLastMsg` already returns.

Why a plain "last seq for subject" isn't enough on its own: `si.sseq` must be the **origin** stream sequence taken from the `Nats-Stream-Source` header, not the hub's storage sequence. So we use the index to *find* the source's last stored message in $O(1)$, then read the origin seq from its header.

3. Background: the subject index (`psim` + `fss`) and the source header

The filestore keeps a **two-level subject index**. Phase 1 stands entirely on it, so it's worth being precise about what each level is.



3.1 `psim` — the store-wide "which blocks hold this subject" map

`fs.psim` (`filestore.go:195`) is a `stree.SubjectTree[psi]` — a subject-token tree keyed by the full message subject — whose value is a tiny per-subject record (`filestore.go:168`):

```
type psi struct {
    total uint64 // how many messages exist for this subject across the whole s
    fblk  uint32 // index of the FIRST block that still contains this subject
    lblk  uint32 // index of the LAST block that still contains this subject
}
```

It is maintained incrementally on every write (`filestore.go:4933-4943`): on store, the subject's entry is created (`total:1, fblk=lblk=current block`) or its `total` is bumped and `lblk` advanced to the current block; on removal/expiry the counts and `fblk` / `lblk` are walked forward as blocks empty out. Because `psim` is a *subject tree*, it supports both exact `Find(subject)` and wildcard `Match(filter)` in time proportional to the subject token structure — **not** to the number of messages.

The single field that makes Phase 1 cheap is `lblk` : given a source's subject, `psim.Find` returns the **exact last block** that contains it. There is no need to scan newer blocks — most of which don't contain that subject at all. (`MultiLastSeqs` , `filestore.go:3806` , uses the same `psim` to answer a *batch* of filters in one index pass.)

3.2 `fss` — the per-block "first/last seq of this subject in this block" map

`psim` gets us to a block; `mb.fss` (`filestore.go:239`) finishes the job *inside* that block. Each `msgBlock` carries its own `stree.SubjectTree[SimpleState]` mapping subject → (`store.go:180`):

```
type SimpleState struct {
    Msgs  uint64 // messages for this subject in THIS block
    First uint64 // first sequence for this subject in this block
    Last  uint64 // last  sequence for this subject in this block
    // (firstNeedsUpdate / lastNeedsUpdate: lazily recomputed when interior dele
}
```

So the two levels compose into an O(1)-ish lookup, which is exactly what `loadLast` (`filestore.go:8980-9034`) does:


```
psim.Find(subject) → info.lblk → block → block.fss.Find(subject) → ss.Last
```

`fss` is what gives us the precise sequence within the block (and it correctly skips interior-deleted messages via `lastNeedsUpdate` / `recalculateForSubj`). On a cold block, `ensurePerSubjectInfoLoaded` loads/derives `fss` for that **one** block — the cost we pay once per resolved source, versus the current scan paying it for every block back to the oldest source.

Two-level traversal: psim subject tree → the block → that block's fss subject tree → the message


1 fs.psim — store-wide subject tree (subject → psi{total, fblk, lblk})

2 follow lblk → the last block holding that subject



3 B2.fss — per-block subject tree (subject → SimpleState{Msgs, First, Last})

4 load msg at Last (hub seq 5123) → read its header



3.3 Why each message carries a `Nats-Stream-Source` header

psim / fss index by **subject**. They know nothing about *which* source produced a stored message or *where in the origin stream* it came from — and those are precisely the two facts resume needs. That information is carried only in the per-message header `Nats-Stream-Source` (`JSStreamSource`, `stream.go:635`), written by `genSourceHeader` (`stream.go:4470`) as:

```
Nats-Stream-Source: <idName> <originSeq> <filter> <destination> <origSubject>
                    |           |
                    |           └ the ORIGIN stream's sequence – what we must
                    └ origin stream name (+ domain/consumer hash); with filter
                        reconstructs the source's full iname (parsed by stream)
```

This header is needed for three independent reasons:

1. **Source attribution.** A hub message is stored under its *destination* subject; the storage layer records nothing about its origin. When two sources can land on the same/overlapping subject, only the header's `iname` says which source a given message belongs to — hence Phase 1's `iname == si.iname` check.
2. **Origin-sequence tracking.** The sourcing consumer is (re)created on the **origin** with `DeliverByStartSequence`, which lives in the *origin*'s sequence space. The hub's own storage sequence (e.g. 5123) is unrelated to the origin sequence (e.g. 482). Only the

header preserves the origin sequence, so it is the *only* place a resume point can be recovered from without contacting the origin.


3. Loop / daisy-chain handling & dedup. When re-sourcing ($A \rightarrow B \rightarrow C$),

`processInboundSourceMsg` strips the inbound `JSSource` and stamps its own (`stream.go:4398`, `4405`), so each hop's provenance is well-defined and re-sourced duplicates can be recognised.

This is the crux of the whole proposal: **the index tells us *which block* in $O(1)$; the header (read from that one message) tells us *which source* and *which origin seq***. Today's scan reads the header off *every* message on the way back because it never consults `psim / fss` to jump directly to the right one.

4. Proposed algorithm

Two phases. Phase 1 resolves the common case with index lookups; Phase 2 keeps today's scan only for sources that genuinely can't be attributed by subject.



For each source, its **stored subject** is the `FilterSubject` for a plain source, or — for a source with a **single, concrete (non-templated)** subject transform — the transform **Destination**.

Templated transform destinations contain `{{...}}` mapping tokens (the rendered subject varies per message), and partial wildcards in a destination are rejected at config time, so a destination is concrete unless it is empty, `>`, or contains `{{`.

Phase 1 – index lookup (per source, $O(1)$ targeted load):
for each source `si`:

```

    subj := si.FilterSubject // plain source
    if si has exactly one transform: subj := transform.Destination
    if subj == "" || subj == ">" || subj contains "{": defer to Phase 2; continue
    sm := store.LoadLastMsg(subj) // jumps via psim, ~1 block
    if sm has a JSStreamSource header:
        (_, iname, sseq) := streamAndSeq(header)
        if iname == si.iname: // self-verifying attribution
            si.sseq = sseq; mark resolved
    // anything not resolved (no header / header for a different source / wildcard
    // multi-or-templated transform / pre-2.10) → Phase 2

```

Phase 2 – fallback reverse scan (only for the unresolved subset):

```

    run today's LoadPrevMsgMulti loop, but with the sublist built from ONLY the
    unresolved sources, so it narrows immediately and stops as soon as they are found

```

The header check makes Phase 1 **self-correcting**: if a subject is shared between sources, or overlaps a direct publish on the hub, or the last message has no source header, the lookup simply doesn't match and that source drops to Phase 2 — never producing a wrong answer.

5. Before → after (code)

Before (`startingSequenceForSources` , condensed — `stream.go:4694`)

```

mset.resetSourceInfo()
// build a sublist of ALL sources' subjects
refreshSublist() // every source
for last := state.LastSeq; ; { // walk the WHOLE store backwards
    sm, seq, err := mset.store.LoadPrevMsgMulti(sl, last, &smv) // loads/decompresses
    if err == ErrStoreEOF || err != nil { break }
    last = seq - 1
    ...
    _, iName, sseq := streamAndSeq(bytesToString(ss))
    update(iName, sseq) // narrows sublist as sources are found
    if len(seqs) == expected { return } // stop only when ALL sources are found
}

```

Cost: O(blocks back to the least-recently-active source) block loads → worst case the entire store.

After

```
mset.resetSourceInfo()
var smv StoreMsg
unresolved := map[string]*sourceInfo{}

// Phase 1: targeted index lookup per source.
for _, ssi := range mset.cfg.Sources {
    si := mset.sources[ssi.iname]
    if si == nil { continue }
    var subj string // FilterSubject, or single c
    switch {
    case len(ssi.SubjectTransforms) == 0: subj = ssi.FilterSubject
    case len(ssi.SubjectTransforms) == 1: subj = ssi.SubjectTransforms[0].Destin
    }
    if subj == _EMPTY_ || subj == fwcs || strings.Contains(subj, "{{") {
        unresolved[ssi.iname] = si // can't attribute by a single
        continue
    }
    sm, err := mset.store.LoadLastMsg(subj, &smv) // 0(1) via psim/fss, returns
    if err != nil || sm == nil || len(sm.hdr) == 0 {
        unresolved[ssi.iname] = si
        continue
    }
    if ss := sliceHeader(JSStreamSource, sm.hdr); len(ss) > 0 {
        if _, iName, sseq := streamAndSeq(bytesToString(ss)); iName == ssi.iname
            si.sseq, si.dseq = sseq, 0 // verified – resolved without
            continue
        }
    }
    unresolved[ssi.iname] = si
}

// Phase 2: only the leftovers pay the (now tiny) reverse scan.
if len(unresolved) > 0 {
    mset.reverseScanForSources(unresolved) // = today's loop, scoped to
}
```

Cost: Phase 1 = $O(S)$ targeted lookups (recent sources share/cache the last block; a quiet source is **one** targeted load, not a walk). Phase 2 only runs for ambiguous sources, with a smaller sublist.

Multi-destination transform sources: call `LoadLastMsg` for each destination and keep the highest origin `sseq`, or defer them to Phase 2. Either is fine; deferring is simplest to start.

6. Before → after (complexity)

	Before	After
Index used	none (generic message walk)	<code>psim</code> / <code>fss</code> per-subject last-block index
Blocks loaded	back to least-active source ($\leq$ all blocks)	~1 per source, only the blocks actually holding a source's last msg
Sensitivity to a quiet source	drags the scan back across the whole store	none — jumps straight to its block
Sensitivity to store size	linear in blocks scanned	none
Worst case still possible?	normal case	only if every source is <code>> /wildcard/shared-subject/pre-2.10</code> (→ Phase 2 == today)
Held lock	<code>mset.mu</code> for the whole walk	<code>mset.mu</code> for O(S) lookups; Phase 2 only for leftovers


The edge → hub topology (each edge mapped to a distinct subject/domain) lands entirely in Phase 1 → the backward scan disappears for that case.

7. Phase 2: keeping the fallback scan narrow

Phase 1 removes most sources from the picture, but whatever remains still runs the reverse `LoadPrevMsgMulti` scan. That scan's whole speed advantage depends on the **sublist being narrow** — `prevMatchingMulti` uses each block's `fss` to skip blocks that contain none of the sublist's subjects. A single `>` (full wildcard) in the sublist defeats that: it matches every subject, so no block can be skipped and the scan visits every block back to the match.

Phase 2 sublist narrowing: one transform source no longer poisons the whole scan

BEFORE — transform/catch-all source injects ">" (empty FilterSubject)



AFTER — phase 1 resolves wildcard transforms; phase 2 sublist uses destinations

step 1: phase 1 collapses wildcard(/)\$N tokens and resolves via the index



step 2: only genuinely-unresolved sources remain; sublist stays narrow



Remaining ">" cases (by design): a genuine catch-all source (empty filter, no transform) or an exotic transform (partition/split).

What used to inject >

1. **Every transform source.** A transform source has an empty `FilterSubject` (mutually exclusive with `SubjectTransforms`, `stream.go:1976`), so the old `refreshSublist` did `s1.Insert(fwcs)`. One transform source in the unresolved set widened the whole sublist to `>` → full backward scan.
2. **Templated transforms were deferred to Phase 2 at all.** Their *rendered* destination is a narrow wildcard (e.g. `tout.{{wildcard(1)}}` always lands on `tout.*`), so they don't need a full scan.
3. **Dead entries.** Phase 2 also inserted `si.sfs` — the transform *source* filters (e.g. `tin.*`) — which never match stored subjects (messages are stored under the *destination*).

The two changes

- **Resolve wildcard transforms in Phase 1.** `transformUntokenize(dest)` collapses `wildcard()` / `$N` mapping tokens to a subject wildcard (`tout.{{wildcard(1)}}` → `tout.*`); `LoadLastMsg` accepts that wildcard and the header- `iname` check still guards attribution. So the common templated transforms resolve in Phase 1 and never reach the scan.
- **Build the Phase 2 sublist from destinations.** For any transform source that *does* reach Phase 2 (e.g. its wildcard space is shared with another source, so Phase 1 saw a header for the other one), insert the destination's wildcard form instead of `>`, and drop the dead `si.sfs` inserts.

What still forces > (by design)

Only a source that *genuinely* needs a broad match:

- a **catch-all source** — empty `FilterSubject`, no transform — sourcing every subject of its origin; and
- an **exotic transform** (`partition` / `split` / `slice` /...) whose rendered destination isn't a subject wildcard, so `transformUntokenize` can't reduce it (it leaves the `{{...}}` token in place, which we detect and fall back to `>`).

Both are uncommon for the edge → hub fan-in. Distinct subjects, wildcard filters, and `wildcard()` / `$N` transforms all either resolve in Phase 1 or keep the Phase 2 sublist narrow enough to skip blocks.

Correctness note: this is purely about *narrowing*, never about *missing* a source. If a narrowed subject is wrong (shared space), the header check fails and the source stays in the unresolved set; an exotic/empty case falls back to `>` . So Phase 2 still finds every source it did before — it just loads far fewer blocks getting there.

8. Correctness & edge cases

- **Same semantics.** Today's scan records, per source, the **most recent** stored message's origin seq (first hit scanning backward). `LoadLastMsg` returns exactly that message. Deleted/interior messages are handled by `loadLast` (it skips `dmap` entries and walks to the previous block if needed), matching the scan's behaviour.
- **Shared subjects / direct publishes.** Resolved safely by the header-iname verification → fall to Phase 2.
- **Wildcard / empty (`>`) filters.** Can't be attributed to one source by subject → Phase 2 (same as today for them).
- **Pre-2.10 headers** (no `iname`, only stream name): verification won't match → Phase 2, which keeps the existing stream-name matching path.
- **memstore.** `MemStore` has the simpler linear `LoadPrevMsgMulti`; `LoadLastMsg` there is also index-light, but memstore streams are small, so Phase 2 cost is negligible. (We can add a memstore `loadLast` fast path later if needed.)
- **setStartingSequenceForSources** (the `STREAM.UPDATE` path, `stream.go:4566`) gets the same Phase 1 treatment for the subset of sources it processes.

9. Risks

- **Index freshness.** `psim` / `fss` may lazily need a recalculation (`lastNeedsUpdate`, `recalculateForSubj`); `loadLast` / `MultiLastSeqs` already handle that, so we inherit

correct behaviour.

- **Behavioural parity.** Phase 1 must produce identical `si.sseq` to the old scan for non-ambiguous sources; this is the core test (below). Keeping Phase 2 byte-for-byte as today bounds the blast radius.
- **Encryption/compression.** Phase 1 still loads the one block holding a source's last message (decrypt/decompress), but only that block — no change in correctness, large reduction in volume.

10. Testing & validation

Prototype status — implemented

A working prototype of the two-phase resolver is implemented in `startingSequenceForSources` (`server/stream.go`), with:

- `TestJetStreamStartingSequenceForSourcesIndexFastPath` (`server/jetstream_sourcing_resume_test.go`) — three distinct-subject sources (index fast path) plus one **subject-transform** source (phase 2 fallback), each with a different origin depth and buried under 20k direct publishes. Asserts every recovered `si.sseq` equals the expected last origin sequence. **Passes.**
- `TestJetStreamStartingSequenceForSourcesAmbiguity` — a catch-all (empty filter) source and **two sources transformed onto the same destination subject** (shared stored subject), mixed with a distinct-subject source. Asserts the phase 2 fallback disambiguates all of them. **Passes.**
- `TestJetStreamSetStartingSequenceForSourcesIndex` — the same fast path applied to the `STREAM.UPDATE` twin `setStartingSequenceForSources` (distinct subject + catch-all + a **concrete-destination transform** source), clearing and recovering `sseq`. **Passes.**
- **Transform fast path:** a source with a single subject transform is resolved in Phase 1 by collapsing `wildcard()` / `$N` mapping tokens to a wildcard form (`transformUntokenize` , e.g. `tout.{{wildcard(1)}}` → `tout.*`) and doing `LoadLastMsg(wildcardForm)` , verified by the header. Only multi-transform and **exotic** transforms (`partition` / `split` / ... that don't reduce to a subject wildcard) still defer to Phase 2.
- **Phase 2 no longer degrades to `>` for transform sources.** Previously a transform source (empty `FilterSubject`) inserted a `>` catch-all into the sublist, defeating per-block skipping for the whole scan. The Phase 2 sublist is now built from the transform **destinations** (wildcard form), with the dead `si.sfs` source-filter inserts removed; only a genuine catch-all source or an exotic transform still forces `>` .
- The change was also applied to `setStartingSequenceForSources` itself (`stream.go:4566`), so both the leader-election and the config-update resume paths use the index fast path and the narrowed sublist.

- Tests (all pass, incl. `-race`): `...IndexFastPath`, `...Ambiguity` (catch-all + shared concrete destination), `...TemplatedTransform` (wildcard transform), `...SharedWildcardTransform` (two templated transforms sharing a `gout.*` space), and `TestJetStreamSetStartingSequenceForSourcesIndex` (twin).
- `BenchmarkJetStreamScanForSources` (existing, single source), `BenchmarkJetStreamScanForSourcesMulti` (16 sources spread across the store), and `BenchmarkJetStreamSourceResumeLeafnodeFanIn` (8–512 edges feeding a hub, quiet edges buried under an all-sourced tail — see the fan-in table below).
- Existing sourcing suite (`SourceBasics`, `SourceRemovalAndReAdd`, `WorkQueueSourceRestart`, `SourceWorkingQueueWithLimit`, `StreamSourceWithoutDuplicateWindow`, `MirrorAndSourcesFilteredConsumers`) still passes.

Measured (filestore, single-server)

Benchmark	Before (reverse scan)	After (phase 1)	Speedup
<code>ScanForSources</code> — 1 source, ~100k buried msgs	~8.5 µs/op	~5.5 µs/op	~1.5×
<code>ScanForSourcesMulti</code> — 16 sources, spread	~110.8 µs/op	~14.0 µs/op	~7.9×

The single-source gap is modest because `LoadPrevMsgMulti` already skips non-matching messages per block; the win scales with the number of sources (each avoids a re-walk via a direct `psim` jump), which is exactly the edge → hub fan-in case.

Leafnode fan-in resume (the leader-election case)

`BenchmarkJetStreamSourceResumeLeafnodeFanIn` models the scenario directly: a hub sources from N edge/leafnode streams, then one chatty edge sources a deep 100k tail that buries the other edges' last messages. Crucially the hub store is **entirely sourced edge subjects** — as a real fan-in is — so every block matches the recovery sublist and *none* can be skipped. This is the work a freshly elected leader runs (`startingSequenceForSources`) before it can resume sourcing. One op == one full recovery.

edges	Before (reverse scan)	After (phase 1)	Speedup	allocs/op before → after
8	10.1 ms	0.27 ms	37×	324 → 46
32	10.8 ms	0.38 ms	28×	3,325 → 173
128	15.8 ms	0.73 ms	22×	44,358 → 657

edges	Before (reverse scan)	After (phase 1)	Speedup	allocs/op before → after
512	101.2 ms	2.00 ms	51×	674,344 → 2,572

(filestore, single-server, `-benchtime=50x`, both columns on the same 2.8 GHz host; "before" measured by splicing in the pre-optimization function from the parent commit.)

Two effects compound, and the table separates them:

- **Store-bound vs source-bound.** "Before" carries a ~10 ms floor *independent of edge count* — it must reverse-scan the whole 100k tail to reach the buried quiet edges. "After" is flat in tail depth and scales only with edge count (one `LoadLastMsg` index jump each).
- **The $O(\text{sources}^2)$ sublist rebuild.** The old Phase 2 rebuilt the whole sublist on *every* source it resolved (`refreshSublist` per `update`). At 512 sources that is ~260k inserts per recovery — the "before" jumps to 101 ms and **49.8 MB / 674k allocs**. Phase 1 resolves these without ever building a sublist, so allocations drop by ~260× at 512 edges.

Caveat: this is a micro-benchmark of the recovery function, not a full Raft election — an actual election adds a fixed quorum/round-trip cost on top of *both* columns equally. An earlier draft buried the edges under *direct* (non-sourced) hub traffic; there the old per-block `fss` skipping rescued the scan (only ~5×), which is not representative of a pure fan-in. The committed benchmark uses an all-sourced tail so nothing is skippable.

What still forces a full `>` scan in Phase 2 (by design)

After the above, Phase 2 only widens to `>` (no block skipping) when a source *genuinely* needs it:

- a **catch-all source** — empty `FilterSubject`, no transform — that sources every subject of its origin (it really can land anywhere, so a broad match is correct); and
- an **exotic transform** (`partition` / `split` / `slice` /...) whose rendered destination is not a subject wildcard, so `transformUntokenize` can't reduce it to a matchable pattern.

Both are uncommon for the edge → hub fan-in. Everything else — distinct subjects, wildcard filters, `wildcard()` / `$N` transforms — either resolves in Phase 1 or narrows the Phase 2 sublist to a concrete or wildcard subject.

Still to do

- **Pre-2.10 / direct-publish-overlap coverage:** add explicit cases (handled today via the header check and the Phase 2 fallback; tests would lock the behaviour in). Hard to construct through the JS client because streams can't declare overlapping subjects; would need low-level store seeding.

- A `partition / split` transform benchmark to confirm the (rare) `>` path is acceptable.

Resolved finding (was suspected pre-existing bug)

An earlier note suspected `setStartingSequenceForSources` mishandled transform sources because its Phase 2 sublist used `si.sfs` (transform **source** filters) rather than the **destination**. On inspection this was **not** a correctness bug: transform sources have an empty `FilterSubject`, so the old sublist inserted a `>` catch-all that matched the destination anyway (verified by `TestJetStreamStartingSequenceForSourcesTemplatedTransform`). It was purely an **efficiency** problem — the `>` defeated block skipping — now fixed by building the Phase 2 sublist from the destination wildcard form and dropping the dead `si.sfs` inserts.

11. Relationship to the durable-consumer proposal

This change is **orthogonal and complementary** to the durable-consumer/ `si.sseq` -persistence ideas in `sourcing-durable-resume.md`:

- Durable consumers / persisted `si.sseq` aim to **skip** resume work entirely (when state survives).
- This proposal makes the **fallback resume itself cheap** — which still runs on cold start, on snapshot restore, for limits/clustered streams, and any time persisted state is missing.

Recommended order: land this index-based resume first (self-contained, no protocol/state change, helps every retention type today), then layer the persistence/durable improvements on top.

Appendix — key references

Symbol	File:line	Role
<code>startingSequenceForSources</code>	<code>stream.go:4694</code>	the scan to replace (Phase 1 + Phase 2)
unconditional call	<code>stream.go:4821</code>	runs on every leader election / restart
<code>setStartingSequenceForSources</code>	<code>stream.go:4566</code>	update-path twin, same treatment
sublist build (stored subjects)	<code>stream.go:4742-4757</code>	defines a source's stored subject(s)

Symbol	File:line	Role
<code>LoadLastMsg</code> / <code>loadLast</code>	<code>filestore.go:9048</code> / <code>8939</code>	index-based last-msg-for-subject (returns header)
<code>MultiLastSeqs</code>	<code>filestore.go:3806</code>	batched last-seq-per-filter via index
<code>psi</code> / <code>fs.psim</code>	<code>filestore.go:168</code> / <code>195</code>	per-subject <code>{total, fblk, lblk}</code> ; store-wide subject → blocks index
<code>mb.fss</code> / <code>SimpleState</code>	<code>filestore.go:239</code> / <code>store.go:180</code>	per-block subject → <code>{Msgs, First, Last}</code> index
<code>psim</code> maintenance on write	<code>filestore.go:4933-4943</code>	how <code>total</code> / <code>lblk</code> are kept current
<code>LoadPrevMsgMulti</code>	<code>filestore.go:9403</code> / <code>memstore.go:1991</code>	the backward walk used by Phase 2
<code>Nats-Stream-Source</code> header	<code>stream.go:635</code>	constant; the per-message source provenance
<code>genSourceHeader</code> / <code>streamAndSeq</code>	<code>stream.go:4470</code> / <code>4545</code>	writes / parses origin stream, iname, origin seq