

BUG FIX · 2.12 LINE

Filestore restores a stale `fb1k` for `MaxMsgsPer>1` with FIFO removal

Why per-subject message lookups could fail after a restart on nats-server 2.12 — and why 2.14 is structurally immune.

PR: [nats-io/nats-server#8285](https://github.com/nats-io/nats-server/pull/8285) File: `server/filestore.go`

Affects: file-backed JetStream, `MaxMsgsPer > 1` Trigger: FIFO eviction + restart

TL;DR

The filestore keeps a per-subject pointer, `fb1k` (“first block”), so it can find a subject’s oldest message without scanning every block on disk. On the **2.12** line, when a subject with several copies is trimmed down to a single message by **FIFO eviction**, `fb1k` can be left pointing at a block that no longer holds any message for that subject. Because of how `index.db` serializes single-message subjects, this **stale pointer gets persisted**, and after a **restart** lookups start scanning from the wrong place and **fail to find the message**. PR #8285 restores a one-line correction (`info.fb1k = info.lb1k`) that fixes the pointer before it is persisted. On **2.14** the on-disk format always stores both pointers, so the bug cannot occur.

Background: what `fb1k` and `lb1k` are for

A file-backed stream is split into many message **blocks** on disk (`1.lb1k`, `2.lb1k`, ...). To answer “what is the first/oldest sequence for subject X?” quickly, the store keeps a small per-subject info record (`psi`) in memory:

```
type psi struct {
    total uint64 // how many messages this subject currently has
    fb1k  uint32 // first block that holds a message for this subject
```

```
lblk uint32 // last block that holds a message for this subject
}
```

Lookups use this pair as a search window instead of scanning the whole store

(`server/filestore.go`, `firstSeqForSubj`):

```
// See if we can optimize where we start.
start, stop := fs.blks[0].index, fs.lmb.index
if info, ok := fs.psim.Find(stringToBytes(subj)); ok {
    start, stop = info.fblk, info.lblk // <-- only scan [fblk .. lblk]
}
```

So correctness hinges on one invariant: **fblk** must point at (or before) the block that actually contains the subject's oldest surviving message. If **fblk** points past it, the message is invisible to the scan.

The key detail: how single-message subjects are persisted

When stream state is written to `index.db`, the 2.12 format does **not** store **lblk** for a subject that has only one message — it stores **total** and **fblk**, and on restore it reconstructs **lblk** = **fblk**:

```
// decode (2.12 behavior)
psi := psi{total: readU64(), fblk: uint32(readU64())}
if psi.total > 1 {
    psi.lblk = uint32(readU64())
} else {
    psi.lblk = psi.fblk // single-msg subject: lblk is *derived* from fblk
}
```

This is the heart of it: for a one-message subject, **fblk** is the only pointer that survives a restart, and **lblk** is rebuilt from it. So if **fblk** is wrong when **total == 1**, then after a restart *both* pointers are wrong, and the search window **[fblk .. lblk]** collapses onto the wrong block.

Why the pointer goes stale: lazy fblk updates

On removal, the store deliberately does **not** recompute fblk right away — it's a lazy optimization, normally corrected later during a lookup (removePerSubject):

```
// We do not update sense of fblk here but will do so when we resolve during lookup.
if info, ok := fs.psim.Find(bsubj); ok {
    info.total--
    // ... fblk left unchanged ...
}
```

That laziness is fine as long as a lookup happens before the state is persisted. The dangerous combination is MaxMsgsPer > 1 (a subject legitimately living across several blocks) trimmed by **FIFO eviction** (oldest-first), so the surviving message ends up in a later block than fblk — and then the state is written out and the server restarts before any lookup corrects it.

Walkthrough

- 1 Subject foo with MaxMsgsPer = 3 has copies in blocks 1, 4, and 7.

psi = {total: 3, fblk: 1, lblk: 7}



red outline = fblk · green outline = lblk

- 2 FIFO eviction removes the two oldest copies (blocks 1 and 4). Only the copy in block 7 survives.

total drops 3 → 1, but fblk is **not** recomputed (lazy).



psi = {total: 1, fblk: 1 (stale!), lblk: 7}

- 3 Stream state is flushed to index.db. Because total == 1, only fblk = 1 is written; lblk is not persisted.

- 4 Server restarts. On restore: fblk = 1 and lblk = fblk = 1.

psi = {total: 1, fblk: 1, lblk: 1} — search window is now just block 1.

5 Lookup for `foo` scans only block 1, finds nothing, and the message in block 7 is effectively lost.

The fix (PR #8285)

The PR restores a block that had been removed in 2.12.7. When FIFO removal brings a subject down to a single message, it eagerly corrects `fblk` to `lblk` — the block where the survivor actually is — so the value that gets persisted is correct:

```
// If we only ever store one/last message for a subject,
// can correct the first block to where we've just written.
if info != nil && info.total == 1 && mmp == 1 {
    info.fblk = info.lblk
}
```

Now at step 2 above, `fblk` becomes 7, the persisted pointer is correct, and after restart the lookup scans block 7 and finds the message.

Test coverage

The PR adds a regression test in `server/filestore_test.go` that builds a subject spanning multiple blocks under `MaxMsgsPer>1`, drives FIFO removal down to one message, restarts the store, and asserts the message is still retrievable (i.e. `fblk` is not stale).

Why 2.14 is no longer affected

The 2.14 line fixes this at the **format level** rather than by patching the removal path. The newer `index.db` encoder **always** writes both `fblk` and `lblk` for every subject, regardless of `total`:

```
// encode (2.14) - always persists both pointers
buf = binary.AppendUvarint(buf, psi.total)
buf = binary.AppendUvarint(buf, uint64(psi.fblk))
buf = binary.AppendUvarint(buf, uint64(psi.lblk)) // always written
```

```
// decode (2.14) – reads lblk whenever the format version supports it
psi := psi{total: readU64(), fblk: uint32(readU64())}
if psi.total > 1 || version >= 4 {
    psi.lblk = uint32(readU64())
} else {
    psi.lblk = psi.fblk
}
```

Two things follow from this:

- `lblk` is no longer *derived* from `fblk` on restore — it's read directly from disk and remains correct (e.g. block 7) even if `fblk` is stale.
- Because `lblk` survives accurately, the lazy-`fblk` design works as intended: after restart the search window `[fblk .. lblk]` still includes the surviving message's block, the lookup finds it, and the same lookup then self-corrects `fblk` back to the right block.

In short: on 2.12 the survivor's location had to be encoded into `fblk` (because `lblk` wasn't persisted for single-message subjects), so a stale `fblk` was fatal. On 2.14 the location is persisted independently in `lblk`, so a temporarily-stale `fblk` is harmless and self-healing. PR #8285 is the backport-shaped workaround that gives 2.12 the same guarantee without changing its on-disk format.

Version comparison

The `info.fblk = info.lblk` correction **originally existed** and was removed in **2.12.7**, which is what introduced the regression. PR #8285 restores it. So releases at or below **2.12.6** were never affected.

Behavior	≤ 2.12.6	2.12.7 – pre-fix	2.12 + PR #8285	2.14+
Persists <code>lblk</code> for single-msg subject	No (derived from <code>fblk</code>)	No (derived from <code>fblk</code>)	No (derived from <code>fblk</code>)	Yes (always)
Corrects <code>fblk</code> before persisting on FIFO trim-to-one	Yes	No (block removed)	Yes (restored)	N/A (not needed)

Behavior	≤ 2.12.6	2.12.7 – pre-fix	2.12 + PR #8285	2.14+
Message retrievable after restart (<code>MaxMsgsPer>1</code> , FIFO)	Yes	No — bug	Yes	Yes
<code>Nats-Expected-Last-Subject-Sequence</code> correct after restart	Yes	No — bug	Yes	Yes

Which stream configurations are affected

Three conditions must all hold for a subject to be exposed to the bug:

1. **File storage** — the bug only surfaces across a restart via `index.db`; memory streams are immune.
2. `MaxMsgsPerSubject > 1` (including unlimited, `-1`) — a single subject must be able to accumulate multiple copies, which can land in different blocks. With `MaxMsgsPerSubject = 1` the existing write-path correction already keeps `fbk` accurate.
3. **FIFO / oldest-first removal** trimming such a subject down to exactly one surviving message that lives in a *later* block. This is driven by any limit-based eviction: `MaxAge` (TTL), `MaxMsgs`, `MaxBytes`, or `Interest / WorkQueue` retention.

Examples that ARE affected

KV bucket with history > 1 and a TTL — the most common real-world trigger. A KV bucket maps to a stream with `MaxMsgsPerSubject = history`; the TTL drives FIFO expiry of old revisions:

```
nats kv add CONFIG --history=10 --ttl=1h --storage=file
// underlying stream: KV_CONFIG
// max_msgs_per_subject = 10, max_age = 1h, storage = file
```

Limits stream with a per-subject cap plus age-based eviction:

```
{
  "name": "EVENTS",
  "storage": "file",
  "subjects": ["events.>"],
  "retention": "limits",
  "max_msgs_per_subject": 5,           // > 1 -> multiple copies per subject
  "max_age": 3600000000000            // 1h TTL -> FIFO expiry
}
```

Per-subject cap plus a total message/byte limit:

```
{
  "name": "ORDERS",
  "storage": "file",
  "subjects": ["orders.*"],
  "max_msgs_per_subject": 3,           // > 1
  "max_msgs": 1000000                 // total cap -> oldest-first eviction
}
```

Examples that are NOT affected

- **Memory storage** (`"storage": "memory"`) — no `index.db` , no restart restore path.
- **`max_msgs_per_subject = 1`** (e.g. a KV bucket with `--history=1`) — the write-path correction (`filestore.go` : `info.total == 1 && mmp == 1`) already keeps `fblk` correct.
- **No eviction** — a stream that only ever grows (no `MaxAge` / `MaxMsgs` / `MaxBytes` , Limits retention) never performs the FIFO trim-to-one that strands the pointer.
- **Any 2.14+ deployment**, and any release $\leq 2.12.6$, regardless of config.

Are KV streams affected?

Yes — KV buckets created with `history > 1` on file storage are squarely in the blast radius, and are the most likely place to hit this bug in practice.

A JetStream KV bucket is just a stream under the hood, named `KV_<bucket>` , where:

- each **key** maps to a **subject** (`$KV_<bucket>.<key>`), and

- the bucket's **history** setting is translated directly into `MaxMsgsPerSubject = history`.

So a bucket with `history = N (N > 1)` keeps up to N revisions per key — i.e. multiple messages per subject, exactly condition 2 of the bug. Revisions are then trimmed **FIFO** (oldest revision first) whenever history is exceeded or a per-bucket TTL (`MaxAge`) expires old revisions — exactly condition 3. When a key is trimmed down to a single surviving revision that happens to live in a later block, its `fbk` goes stale and is persisted; after a server restart the key's scan window collapses onto the wrong block.

```
// Affected: history > 1 on file storage
nats kv add SESSIONS --history=5 --storage=file
// -> stream KV_SESSIONS, max_msgs_per_subject = 5

// NOT affected: default history = 1 (max_msgs_per_subject = 1)
nats kv add FLAGS --storage=file
```

How KV operations break after a restart

Because every KV read and conditional write resolves the key's *last* revision through the same `LoadLastMsg` path that depends on `fbk` / `lblk`:

- `kv.Get(key)` / `kv.History(key)` — the surviving value sits in a block outside the stale scan window, so the lookup reports the **key as missing** even though the data is intact on disk. Reads appear to lose data until something forces the pointer to self-correct.
- `kv.Update(key, val, rev)` (**compare-and-set**) — KV implements CAS using the `Nats-Expected-Last-Subject-Sequence` header (see below). With the stale pointer the server reports the last revision as `0`, so a correct `Update` at the real revision is **wrongly rejected** with a wrong-last-sequence error.
- `kv.Create(key, val)` (**create-only**) — `Create` guards with expected revision `0` ("key must not exist"). Since the server also believes the key is gone, the create is **wrongly accepted**, silently overwriting / duplicating an existing key and breaking the create-once guarantee.

In short: a KV bucket with `history > 1` on disk can, after a restart, appear to drop keys and make CAS operations return wrong results — even though no data was actually lost. The mechanism for the CAS failures is the expected-last-subject-sequence check, explained next.

Why `Nats-Expected-Last-Subject-Sequence` is affected

This publish header implements per-subject optimistic concurrency: “only accept my message if the last sequence currently stored for this subject equals N .” The server evaluates it in `processJetStreamMsg` (`server/stream.go`) by asking the store for the subject’s last message:

```
// stream.go - expected last sequence per subject
if seq, exists := getExpectedLastSeqPerSubject(hdr); exists {
    sm, err := store.LoadLastMsg(seqSubj, &smv)
    if sm != nil {
        fseq = sm.seq
    }
    if err == ErrStoreMsgNotFound && seq == 0 {
        fseq, err = 0, nil
    }
    if err != nil || fseq != seq { // mismatch -> reject the publish
        resp.Error = NewJSStreamWrongLastSequenceError(fseq)
        ...
    }
}
```

The catch is *how* `LoadLastMsg` finds that last message. For a literal subject it calls `loadLast`, which walks blocks **backwards** using the very same per-subject pointers (`filestore.go`):

```
if info, ok := fs.psim.Find(stringToBytes(subj)); ok {
    start, stop = info.lblk, info.fblk // scan window [lblk .. fblk], backwards
}
// Walk blocks backwards looking for the subject's last message.
for i := start; i >= stop; i-- { /* ... */ }
```

The failure

After a restart in the buggy scenario, the single-message subject restores to `{fblk: 1, lblk: 1}` (the stale value) even though the real message is in block 7. `loadLast` therefore scans only block 1, finds nothing, and returns `ErrStoreMsgNotFound` → `fseq = 0`. The server now believes the subject is empty, which corrupts the header check in **both directions**:

- **Wrongful rejection.** A correct publisher sends `Nats-Expected-Last-Subject-Sequence: 7` (the true last). The server compares `fseq (0) != 7` and rejects with `JSStreamWrongLastSequenceError` — a valid write is refused and the publisher is told the wrong current sequence (`0`).

- **Wrongful acceptance.** A publisher guarding “only if this subject is brand new” sends `Nats-Expected-Last-Subject-Sequence: 0`. Because the server also thinks the subject is empty (`fseq == 0`), the check passes and the message is accepted — silently defeating the optimistic-concurrency guarantee and clobbering/duplicating an existing subject.

So the stale `fblk` isn't just a “message not found on `get`” problem — it poisons every per-subject sequence lookup that rides on those pointers, and the expected-last-subject-sequence header is the most visible casualty because it turns a storage glitch into incorrect publish accept/reject decisions. Restoring `info.fblk = info.lblk` (or, on 2.14, persisting `lbk` directly) keeps the scan window pointing at the real block, so `LoadLastMsg` returns the correct sequence again.

Recovery: rebuild `index.db`

If you are already running an affected version and a stream/KV bucket has stranded a key, you don't need the patched binary to get your data back. The stale pointers live **only** in the per-stream state snapshot, `index.db`. The actual messages are in the immutable block files (`1.blk`, `2.blk`, ...), which are untouched by this bug. Deleting `index.db` forces the server to **rebuild the per-subject state by scanning the blocks** on restart, which recomputes `fblk` / `lbk` from the real data — correctly.

This is safe and non-destructive. `index.db` is a cache/optimization that is always reconstructable from the block files. The server already rebuilds it automatically whenever it detects the snapshot is stale (`filestore.go`: “*Detected a stale index.db ... will rebuild*”). You are just forcing that rebuild.

Procedure (single server)

- 1 **Stop the nats-server.** This is required — the server rewrites `index.db` on clean shutdown, so deleting it on a live server would simply be overwritten (or recreated) with the same stale data.
- 2 **Locate the stream's message directory.** The layout under the JetStream store dir is:

```
<store_dir>/jetstream/<account>/streams/<stream>/msgs/index.db
```

```
# e.g. a KV bucket "SESSIONS" -> stream "KV_SESSIONS":  
/data/jetstream/jetstream/$G/streams/KV_SESSIONS/msgs/index.db
```

(`$G` is the global account; a named account uses its account ID. Confirm your store dir via the server config `jetstream { store_dir }`.)

- 3 **Delete only** `index.db` for the affected stream. Leave every `*.blk` file (and `thw.db` / `sched.db`) in place — those hold your messages.

```
rm /data/jetstream/jetstream/$G/streams/KV_SESSIONS/msgs/index.db
```

- 4 **Restart the nats-server.** On startup it finds no snapshot, walks the block files, and rebuilds `fblk` / `lblk` from the messages that actually exist. The previously “missing” key/message is visible again and the `Nats-Expected-Last-Subject-Sequence` check returns the correct value.

- 5 **Verify**, e.g. `nats stream info <stream>`, a `kv.Get` / `get` on the affected key, or a conditional publish at the expected revision.

Clustered / replicated streams (R>1)

The bug and its on-disk state are **per node**, so the stale `index.db` may exist on one, some, or all replicas. Recover **one node at a time**: stop a single server, delete that node's `index.db` for the stream, restart it, and let it rejoin and catch up before moving to the next. Do not delete on multiple replicas simultaneously — keep a quorum healthy throughout. (An alternative on a healthy cluster is to step down / remove and re-add the affected replica so it re-syncs from the leader.)

Cost & caveats: the rebuild scans all blocks for the stream, so for very large streams expect extra startup time and disk I/O proportional to the stream size. Take a backup/snapshot first if you want belt-and-suspenders. This is a one-shot remediation for

already-affected data; to *prevent* recurrence, upgrade to a build that includes PR #8285 (2.12 line) or to 2.14+.

Scope & limitations

- **2.12 only.** The fix targets the 2.12 release line; 2.14+ already handles this via the format change.
- **FIFO removal path only.** The correction is applied on the limits/eviction path. Other removal patterns on the 2.12 line (e.g. an explicit `RemoveMsg` on the newest message rather than oldest-first eviction) can still leave a stale pointer — a known, documented constraint of 2.12.
- **Low risk.** It restores previously-shipped behavior plus a regression test; no on-disk format change.

Generated explainer for [nats-io/nats-server#8285](https://github.com/nats-io/nats-server/pull/8285). Code references: `server/filestore.go` (`firstSeqForSubj`, `removePerSubject`, `loadLast`, `index.db encode/decode`) and `server/stream.go` (`expected-last-subject-sequence check`, `LoadLastMsg`). For authoritative behavior, consult the PR and the source.